

## 1. Problem Description

The company I work for heavily relies on fish resources as its major raw material. The fish are transported from all over the world to the United States. There are various types of costs associated with transportation and storage, such as demurrage, drayage, cold storage, material handling labor costs, and transportation costs to the factory. All of these costs are classified as variable overhead costs under the group named "fish handling cost," which accounts for approximately 10% of the overall conversion cost for the product. Therefore, an accurate fish handling cash flow forecast is crucial for the business operation and finance.

Given the importance of accurate forecasting, this is a perfect business use case for supervised machine learning technology. Below are the steps to implement this ML solution for fish handling cash forecasting.

## 2. . Data Discovery and Preparation:

### Data Engineering and Architecture.

The volume of fish procured by our company is driven by production volume and schedule, and all related transactions are executed and maintained in our ERP system through the purchase order modules. Once the fish leaves its supplier's port, ownership is transferred to our company, and the transaction is marked as a "good issue" in the ERP system. Upon receipt of the fish at one of our cold storage locations, the transaction is marked as a "good receipt." All of these transactions are stored in the Material Movements tables in the ERP system's database, which drives the fish handling cost.

To determine the amount of cold storage fees charged each month, we must consider the inventory level. Therefore, historical beginning and ending inventory balance data must be collected. Additionally, knowing the production volume from previous periods can help with the ML modeling, as mentioned above.

Finally, the fish handling invoices data can provide insight into the cash being paid each month. With all of this data, we can create a reliable and accurate fish handling cash flow forecast.

### Data Pipelines in Datawarehouse

1. **Inventory Balance Snapshot:** Every day, our ERP system's inventory management module captures a snapshot of the inventory balance. To bring this valuable data into our data warehouse, we have created a data pipeline that extracts the daily inventory snapshot and loads it into the warehouse.
2. **Fish Inbound/Outbound Movements:** The procurement of fish in our organization is facilitated through purchase orders. All inbound and outbound fish movements are recorded in the ERP system and stored in related tables. These tables have been identified in the ERP database and merged into a few tables in our data warehouse through a data pipeline.
3. **Fish Invoice and Payments:** All invoices received against fish receipts in our ERP system are recorded, and the payment transactions are associated with each invoice. These financial records have been extracted from the ERP system and loaded into our data warehouse using a data pipeline.
4. **Production Volume:** Every day, our ERP system stores production records at the production order level. To analyze this data in our data warehouse, we have developed a data pipeline that extracts all production records, aggregates them into monthly buckets, and loads them into a single table in the data warehouse. This makes it easier to analyze production volume trends over time.

## EDA and Data Mining

### 1. Inventory Balance Data: Transformation is done in SQL

```
Select
i.[Material]
-- ,[Material Description]
-- ,try_cast([Plant] as nvarchar) as Plant
-- ,[Storage Location]
-- ,[Batch Number]
-- ,[Material Type]
,SnapshotDate
-- ,try_cast([Stock InQuality SU] as DECIMAL)
-- ,try_cast([InTransit Stock SU] as DECIMAL)
-- ,try_cast([Blocked Stock SU] as DECIMAL)
,Sum(try_cast([Total Stock SU] as DECIMAL)) as TotalQty
-- ,try_cast([RestrictedStock SU] as DECIMAL)
-- ,try_cast([Unrestricted Stock SU] as DECIMAL)
-- ,try_cast([Available] as DECIMAL)
-- ,try_cast([MAC Price] as DECIMAL)
,[UOM] as UoM
-- ,m.MATL_GROUP
-- ,mg.MaterialGroupName as Mat_Grp_Desc
,d.FiscalPeriod
,d.FiscalPeriodWeek
from [SFS].[InventorySnapshots] i
Left JOIN [SC].[MaterialMaster] m
ON i.Material = m.MATERIAL
Left JOIN [GEN].[DimMaterialGroup] mg
ON m.MATL_GROUP = mg.MaterialGroupID
Left Join [GEN].[DimDate] d
ON d.[FullDate]=i.SnapshotDate
Where m.MATL_GROUP in ('NF17','NF28') and Year([SnapshotDate])>=Year(Getdate()-1)-3 and [Storage Location]='WHSE'
Group By d.FiscalPeriod,i.Material
,d.FiscalPeriodWeek,SnapshotDate,UOM
```

Features are the fish material#, Fish material Description, Material group , snapshot date and fiscal Period.

Based on the snapshot date, the Inventory with min of snapshot date in the same fiscal period is labeled beginning inventory for the fiscal period. the Inventory with max of snapshot date in the same fiscal period is labeled ending inventory for the fiscal period.

**Load fish inventory data into Panda dataframe, total the quantity by snapshot date then convert the unit of measure to metric ton.**

```
In [ ]: df = pd.read_excel(r'C:\Users\ilin\Downloads\OFC_1_FishInventory.xlsx')
df=df.groupby(['FiscalPeriod','FiscalPeriodWeek','SnapshotDate'],as_index=False)['TotalQty'].sum()
# df=df.to_frame() #turn searise data into DF

df=df.groupby(['FiscalPeriod','FiscalPeriodWeek'],as_index=False)['TotalQty'].mean()
df['TtlMT']=df['TotalQty']/2204.6
df.tail()
# print(df['FiscalPeriod'])
```

```
5]:
```

|     | FiscalPeriod | FiscalPeriodWeek | TotalQty   | TtlMT       |
|-----|--------------|------------------|------------|-------------|
| 155 | 2023001      | 3                | 10099255.0 | 4580.992017 |
| 156 | 2023001      | 4                | 10170984.0 | 4613.528078 |
| 157 | 2023002      | 1                | 10146631.0 | 4602.481629 |
| 158 | 2023002      | 2                | 10713701.0 | 4859.702894 |
| 159 | 2023002      | 3                | 10835095.0 | 4914.766851 |

**Transform the inventory data to weeks pivottable layout. this makes the beginning and ending inventory logic easy to implement.**

```
df= df.pivot(index=['FiscalPeriod'],columns='FiscalPeriodWeek',values='TtlMT')
# df= df.mean()
```

```
df.head()
```

[7]:

| FiscalPeriodWeek | 1           | 2           | 3           | 4           | 5           |
|------------------|-------------|-------------|-------------|-------------|-------------|
| FiscalPeriod     |             |             |             |             |             |
| 2020001          | 3049.027488 | 3536.496417 | 3103.600200 | 3121.676495 | NaN         |
| 2020002          | 3119.687925 | 2695.267169 | 2359.744625 | 2057.421301 | NaN         |
| 2020003          | 2316.350812 | 2441.844779 | 2053.261816 | 2105.208201 | 2035.600109 |
| 2020004          | 2542.743355 | NaN         | 2694.276966 | 2708.263631 | NaN         |
| 2020005          | 2941.140797 | 2150.299828 | 2207.898485 | 2367.375941 | NaN         |

**Beginning and Ending Inventory Calculation. Column 1 is beginning inventory, and column 4 is the ending, if column 5 is empty.**

```
df['BegInv']=df[1]
df['EndInv'] = df.apply(lambda x: x[5] if x[5] > 0 else x[4], axis=1)
df_Inv=df
df_Inv.head()
```

[8]:

| FiscalPeriodWeek | 1           | 2           | 3           | 4           | 5           | BegInv      | EndInv      |
|------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| FiscalPeriod     |             |             |             |             |             |             |             |
| 2020001          | 3049.027488 | 3536.496417 | 3103.600200 | 3121.676495 | NaN         | 3049.027488 | 3121.676495 |
| 2020002          | 3119.687925 | 2695.267169 | 2359.744625 | 2057.421301 | NaN         | 3119.687925 | 2057.421301 |
| 2020003          | 2316.350812 | 2441.844779 | 2053.261816 | 2105.208201 | 2035.600109 | 2316.350812 | 2035.600109 |
| 2020004          | 2542.743355 | NaN         | 2694.276966 | 2708.263631 | NaN         | 2542.743355 | 2708.263631 |
| 2020005          | 2941.140797 | 2150.299828 | 2207.898485 | 2367.375941 | NaN         | 2941.140797 | 2367.375941 |

## 2. Fish Inbound/Outbound Movement data:

```
SELECT a.[Material]
      ,[Posting Date]
      --,[Plant]
      ,[Fiscal Period]
      ,[Fiscal Year]
      --,[Fiscal Variant]
      --,[Material Document Number]
      --,[Material Document Item]
      --,[Customer]
      ,[Batch]
      --,[Vendor]
      ,[Movement Type]
      ,[Storage Location]
      ,[PO Number]
      ,[Quantity]
      ,[Unit of Measure]
      --,[Sales Order Number]
      --,[Sales Order Item]
      --,[Currency]
      --,[Amount]
      --,[GL Account]
      --,[Chart of Accounts]
      --,[Company Code]
      --,[Reference Document Number]
      --,[Cost Center]
      --,[Controlling Area]
      --,[Long Description]
      --,[Order Number]
      ,[Name]
      --,[b.EQ]
FROM [SC].[MaterialMovements] a
INNER JOIN (Select MATERIAL From [SC].[MaterialMaster]
           Where [MATL_GROUP] in ('NF17','NF28')) b
ON a.[Material]= b.MATERIAL--(REPLICATE('0',18-Len(b.MaterialId))+b.MaterialId)
Where Plant=1100 and [Fiscal Year]>=Year(Getdate())-3
and [Movement Type] in ('351','352') and [PO Number] is not null
```

Features are [Material] ,[Posting Date],[Fiscal Period],[Fiscal Year] ,[Batch] ,[Movement Type],[Storage Location] ,[PO Number] ,[Quantity] ,[Unit of Measure]

Data was pulled with filters of specific material movement types, 351 and 352; and specific material groups, NF17 and NF28.

**Load fish movements data that is downloaded from the data warehouse.**

```
df_Fishmove = pd.read_excel(r'C:\Users\ilin\Downloads\OFC_5_LoinMoves.xlsx')
df_Fishmove=df_Fishmove[['Fiscal Period','Quantity','Unit of Measure']]
df_Fishmove
```

]:

|        | Fiscal Period | Quantity  | Unit of Measure |
|--------|---------------|-----------|-----------------|
| 0      | 2021007       | -2731.499 | LB              |
| 1      | 2021007       | -2821.888 | LB              |
| 2      | 2021007       | -2696.226 | LB              |
| 3      | 2020001       | -1.000    | LB              |
| 4      | 2020005       | -2563.950 | LB              |
| ...    | ...           | ...       | ...             |
| 108040 | 2022012       | 2623.474  | LB              |
| 108041 | 2022012       | 2623.474  | LB              |
| 108042 | 2022012       | 2623.474  | LB              |
| 108043 | 2022012       | 2623.474  | LB              |
| 108044 | 2022012       | 2623.474  | LB              |

108045 rows x 3 columns



**Fish movement is aggregated to fiscal period level in metric tons.**

```
df_Fishmove['Quantity'] = df_Fishmove.apply(lambda x: x['Quantity']/2204.6 if x['Unit of Measure'] == 'LB' else x['Quantity'],
df_Fishmove=df_Fishmove.groupby(['Fiscal Period'],as_index=False)['Quantity'].sum()
df_Fishmove.head()
```

```
[9]:
```

|   | Fiscal Period | Quantity    |
|---|---------------|-------------|
| 0 | 2020001       | 1157.993992 |
| 1 | 2020002       | 1868.092002 |
| 2 | 2020003       | 1617.028000 |
| 3 | 2020004       | 1440.304001 |
| 4 | 2020005       | 1340.953000 |

### 3. Fish Invoice and Payments data:

```
SELECT [InvoiceNumber]
,[InvoiceItem]
,[Material]
,[Plant]
,[Vendor]
,[MaterialType]
,[PostingDate]
,[FiscalYearPeriod]
,[Amount]
FROM [SFS].[PurchaseOrderInvoices]
Where [MaterialGroup] in ('NF17','NF28')
and ClearingDocument='' and GoodsReceiptDate<>'00000000'
```

Features are [InvoiceNumber], [InvoiceItem], [Material], [Plant], [Vendor], [MaterialType], [PostingDate], [FiscalYearPeriod], [Amount] with filters of and specific material groups, NF17 and NF28

**Load Invoice and payment data that was downloaded from the data warehouse. Three years' data was saved into three excel files.**

```
files = [r'C:\Users\ilin\Downloads\OFC_2_2022Payment.xlsx',
r'C:\Users\ilin\Downloads\OFC_3_2021Payment.xlsx',
r'C:\Users\ilin\Downloads\OFC_4_2023Payment.xlsx']

df_Pmt = pd.concat([pd.read_excel(f) for f in files])
df_Pmt = df_Pmt.dropna(subset=['Text'])
df_Pmt = df_Pmt[df_Pmt['Text'].str.contains("Accrual") & (df_Pmt['Amount in local cur.'] > 0)]
df_Pmt.head()
```

```
[9]:
```

|     | DocumentNo | Account | CoCd | Profit Ctr | Vendor Name | User name | Text           | Reference | Pstng Date | Amount in local cur. |
|-----|------------|---------|------|------------|-------------|-----------|----------------|-----------|------------|----------------------|
| 122 | 102456920  | 540474  | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P1 | 840111    | 2022-01-29 | 164004.46            |
| 231 | 102461654  | 540474  | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P2 | 840210    | 2022-02-25 | 219597.91            |
| 295 | 102484392  | 540474  | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P3 | 840314    | 2022-04-01 | 283188.20            |
| 413 | 102486705  | 540474  | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P4 | 840406    | 2022-04-29 | 466399.35            |
| 513 | 102492506  | 540474  | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P5 | 840507    | 2022-05-27 | 595096.53            |

```
Cal = pd.read_excel(r'C:\Users\ilin\Downloads\MasterCalendar.xlsx')
```

```
df_Pmt = df_Pmt.merge(Cal, left_on='Pstng Date', right_on='Date', how='left')
df_Pmt.head()
```

| t | CoCd | Profit Ctr | Vendor Name | User name | Text           | Reference | Pstng Date | Amount in local cur. | Date       | WeekNumber | Quarter | Period | PeriodWeek | FiscPd  |
|---|------|------------|-------------|-----------|----------------|-----------|------------|----------------------|------------|------------|---------|--------|------------|---------|
| 0 | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P1 | 840111    | 2022-01-29 | 164004.46            | 2022-01-29 | 4          | 1       | 1      | 4          | 2022001 |
| 1 | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P2 | 840210    | 2022-02-25 | 219597.91            | 2022-02-25 | 8          | 1       | 2      | 4          | 2022002 |
| 2 | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P3 | 840314    | 2022-04-01 | 283188.20            | 2022-04-01 | 13         | 1       | 3      | 5          | 2022003 |
| 3 | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P4 | 840406    | 2022-04-29 | 466399.35            | 2022-04-29 | 17         | 2       | 4      | 4          | 2022004 |
| 4 | 1001 | 11000      | NaN         | SGOMEZ    | OFC Accrual P5 | 840507    | 2022-05-27 | 595096.53            | 2022-05-27 | 21         | 2       | 5      | 4          | 2022005 |

```
df_Pmt=df_Pmt[['FiscPd','Amount in local cur.']]
df_Pmt.head()
```

|   | FiscPd  | Amount in local cur. |
|---|---------|----------------------|
| 0 | 2022001 | 164004.46            |
| 1 | 2022002 | 219597.91            |
| 2 | 2022003 | 283188.20            |
| 3 | 2022004 | 466399.35            |
| 4 | 2022005 | 595096.53            |

#### 4. Production Volume:

Data is directly pulled from an existing Power BI dashboard which is using SQL pool from our data warehouse. It is leveraging the existing data transformation procedures, to avoid duplicate efforts.

Features are Fiscal Period, and Actual Qty

#### Load production volume data into data frame

```
df_Prd = pd.read_excel(r'C:\Users\ilin\Downloads\OFC_6_Production.xlsx')
df_Prd.head()
```

```
5]:
```

|   | Row Labels | ActStdCases |
|---|------------|-------------|
| 0 | 2020001    | 656337      |
| 1 | 2020002    | 641157      |
| 2 | 2020003    | 861286      |
| 3 | 2020004    | 739870      |
| 4 | 2020005    | 720804      |

### 3. Machine Learning Data Modeling

After data processing steps and EDA, all the datasets are merged into one flat table.

**Final dataset- merge all the datasets from above into one flat table**

```
df_Inv = df.reset_index()
df = df_Inv[['FiscalPeriod', 'BegInv', 'EndInv']].set_index('FiscalPeriod')

df=df.merge(df_Pmt, left_on='FiscalPeriod', right_on='FiscPd', how='inner')
df=df.merge(df_Fishmove, left_on='FiscPd', right_on='Fiscal Period', how='left')
df=df.merge(df_Prd, left_on='FiscPd', right_on='Row Labels', how='left')
df=df[['Fiscal Period', 'BegInv', 'EndInv', 'Amount in local cur.', 'Quantity', 'ActStdCases']]
df.columns=['Fiscal Period', 'BegInv', 'EndInv', 'Payment', 'MoveQty', 'StdCases']
df
```

```
16]:
```

|   | Fiscal Period | BegInv      | EndInv      | Payment     | MoveQty     | StdCases |
|---|---------------|-------------|-------------|-------------|-------------|----------|
| 0 | 2020001       | 3049.027488 | 3121.676495 | 248997.6875 | 1157.993992 | 656337   |
| 1 | 2020002       | 3119.687925 | 2057.421301 | 195793.5625 | 1868.092002 | 641157   |
| 2 | 2020003       | 2316.350812 | 2035.600109 | 196274.8750 | 1617.028000 | 861286   |

Analysis steps:

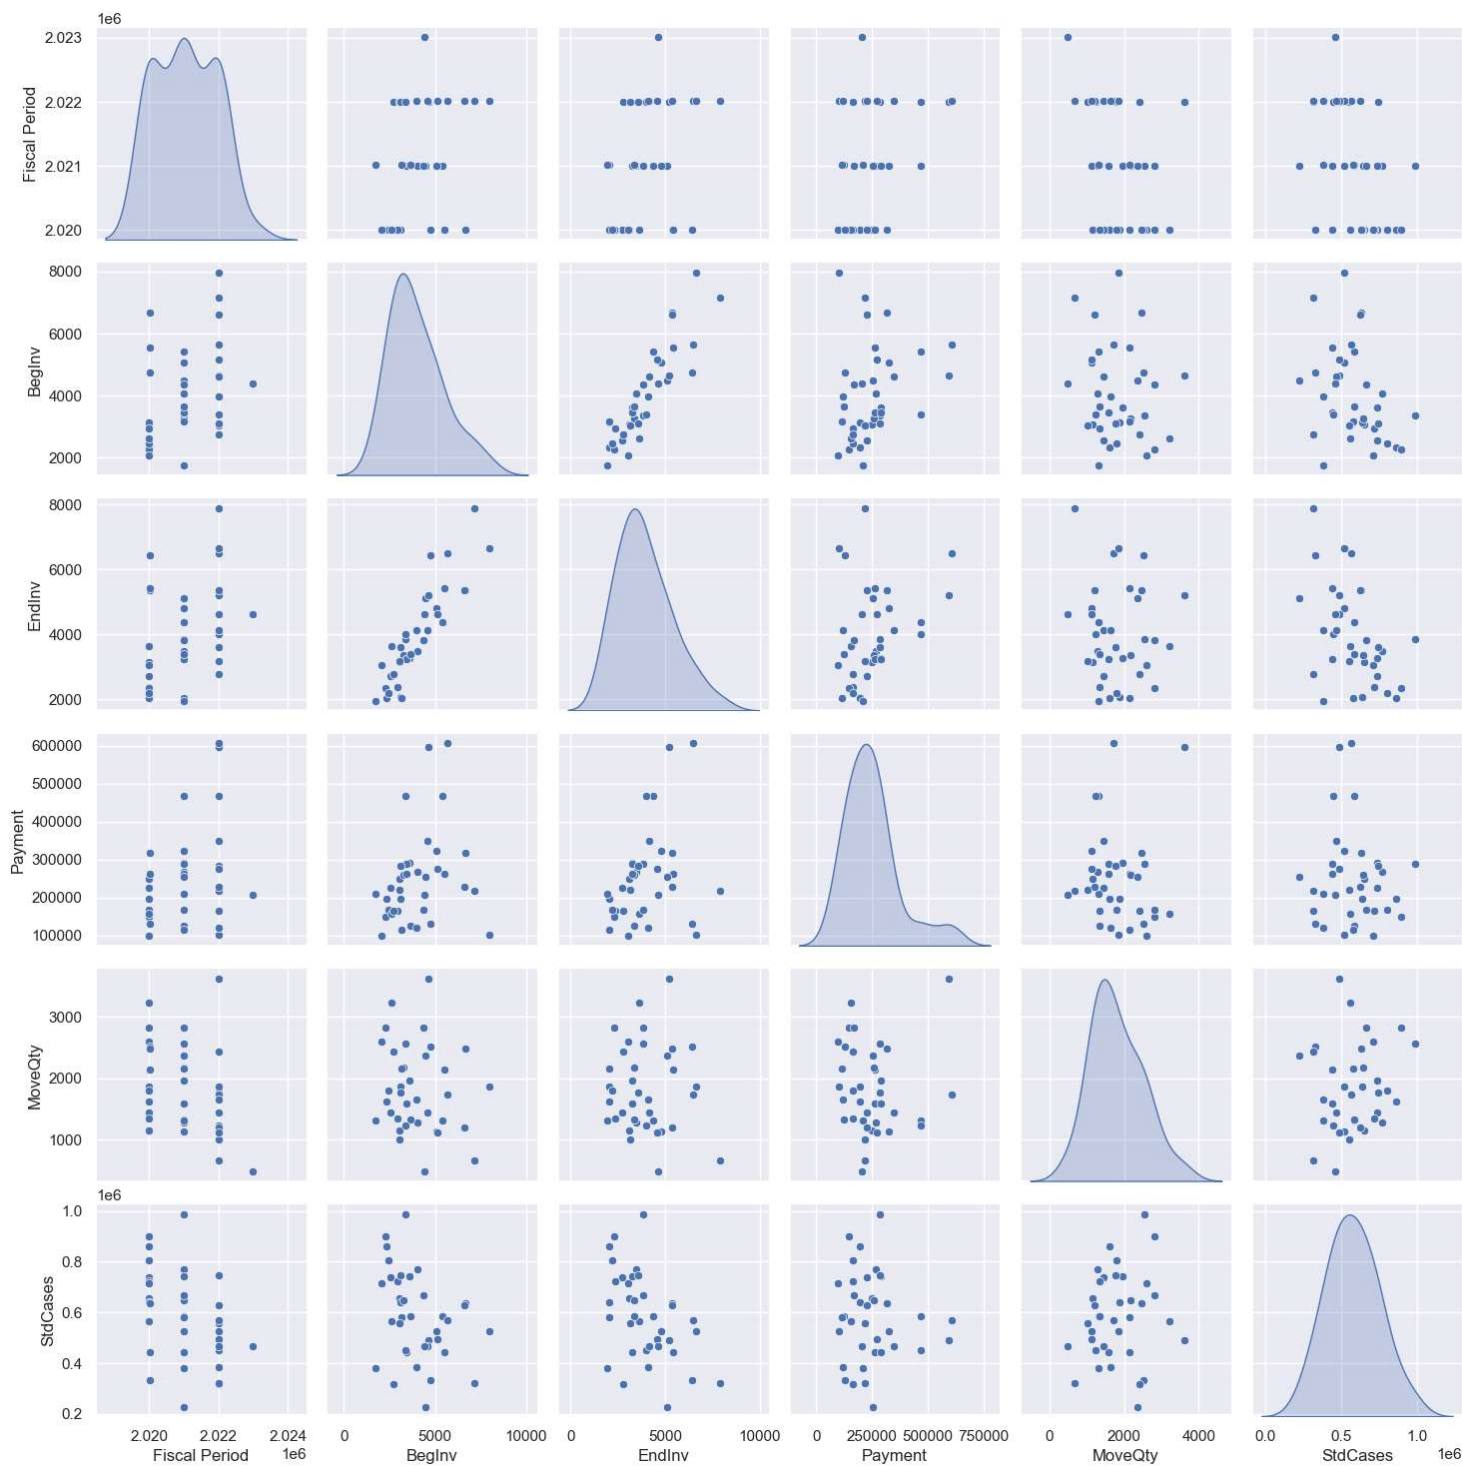
1. Multilinear Regression: since this dataset has multiple independent variables (Beginning Inventory, Ending Inventory, Production cases, Movement Qty) and a single dependent variable (payment)
2. If the accuracy is low and the data shows non-linear relationships with all the variables , then will do polynomial regression.
3. If Polynomial is not ideal, will do random forest regression. Random forest regression is robust to overfitting, which occurs when a model is too complex and fits the noise in the data instead of the underlying signal. This is because the algorithm builds multiple trees, each of which is trained on a random subset of the data, which helps to reduce overfitting.
4. Lastly, will do an AdaBoost regression. The AdaBoost regression algorithm works by building a sequence of weak regression models, each of which is trained on a subset of the data and assigned a weight based on its performance. The algorithm then combines the predictions of all the models to make a final prediction. The weights of the models are adjusted based on the error of the previous model, so that the subsequent models focus more on the data points that were incorrectly predicted by the previous models.

In the end, the final method will be based on the accuracy score from each technics from the above.

First, started with the testing the correlation between each variables. The correlation matrix is plotted in the chart below.

```
df.corr()
sns.pairplot(df, diag_kind = 'kde')
plt.savefig('pair_plot.png', dpi = 300, bbox_inches = 'tight')
```

Correlation Matrix



Run multilinear regression model



```
model= smf.ols(formula= 'Payment~BegInv + EndInv + MoveQty+StdCases ', data=df).fit()
model.summary()
```

18]: OLS Regression Results

|                   |                  |                     |          |       |           |          |
|-------------------|------------------|---------------------|----------|-------|-----------|----------|
| Dep. Variable:    | Payment          | R-squared:          | 0.109    |       |           |          |
| Model:            | OLS              | Adj. R-squared:     | 0.001    |       |           |          |
| Method:           | Least Squares    | F-statistic:        | 1.012    |       |           |          |
| Date:             | Sat, 25 Feb 2023 | Prob (F-statistic): | 0.415    |       |           |          |
| Time:             | 00:15:57         | Log-Likelihood:     | -495.77  |       |           |          |
| No. Observations: | 38               | AIC:                | 1002.    |       |           |          |
| Df Residuals:     | 33               | BIC:                | 1010.    |       |           |          |
| Df Model:         | 4                |                     |          |       |           |          |
| Covariance Type:  | nonrobust        |                     |          |       |           |          |
|                   | coef             | std err             | t        | P> t  | [0.025    | 0.975]   |
| Intercept         | 9.862e+04        | 1.25e+05            | 0.791    | 0.435 | -1.55e+05 | 3.52e+05 |
| BegInv            | -8.9613          | 30.164              | -0.297   | 0.768 | -70.330   | 52.407   |
| EndInv            | 38.3806          | 32.533              | 1.180    | 0.247 | -27.808   | 104.570  |
| MoveQty           | -2.0636          | 29.913              | -0.069   | 0.945 | -62.923   | 58.795   |
| StdCases          | 0.0679           | 0.133               | 0.510    | 0.613 | -0.203    | 0.339    |
| Omnibus:          | 7.609            | Durbin-Watson:      | 1.122    |       |           |          |
| Prob(Omnibus):    | 0.022            | Jarque-Bera (JB):   | 6.266    |       |           |          |
| Skew:             | 0.846            | Prob(JB):           | 0.0436   |       |           |          |
| Kurtosis:         | 4.046            | Cond. No.           | 3.85e+06 |       |           |          |

Based on the correlation matrix and the results from the multi-linear model, it shows the variables are not in a linear relationship. The r-squared score is only 0.109. that suggests we should run polynomial regression.

Run Polynomial regression:

Build an iteration to run through 20 degrees.

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics

X=df[['BegInv','EndInv','MoveQty','StdCases']]
y=df['Payment']
# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
r2=[]
l=[]
for i in range(1,20):
    poly = PolynomialFeatures(degree=i)
    X_poly = poly.fit_transform(X)

    # Fit linear regression model
    reg = LinearRegression().fit(X_poly, y)
    y_pred = reg.predict(X_poly)
    r2.append(metrics.r2_score(y, y_pred))
    l.append(i)
# print("Accuracy of the model:", metrics.r2_score(y, y_pred))


```

```

print(max(r2))
fig, ax = plt.subplots()
ax.plot(l,r2)
# plt.plot(l,r2)
plt.xlabel('Polynomial Degree')
plt.ylabel('R-squared')
plt.title('Relationship between Polynomial Degree and R-squared')
ax.set_xticks(np.arange(0, 20, 1))
plt.show()

```

0.9994176077541682



The best R-squared score is from the 4<sup>th</sup> degree, based on the plot chart, it is apparently overfitted. In this case, the 3<sup>rd</sup> degree is chosen for the model. It has an accuracy score of 0.774

```
poly = PolynomialFeatures(degree=3)
X_poly = poly.fit_transform(X)

# Fit linear regression model
reg_yes = LinearRegression().fit(X_poly, y)
y_pred = reg_yes.predict(X_poly)
print("Accuracy of the model:", metrics.r2_score(y, y_pred))

Accuracy of the model: 0.7743487239098648
```

In order to further improve the accuracy, a random forest model is constructed.

```
from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor()

rf_train=rf.fit(X, y)
rf_pred=rf.predict(X)
print("Accuracy of the model:", metrics.r2_score(y, rf_pred))

Accuracy of the model: 0.8261013602731134
```

The accuracy score is improved to 0.826

Lastly constructed an Adaboost regression model

```
from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor()

# fit the model on the training data
ada.fit(X,y)
ada_pred=ada.predict(X)
print("Accuracy of the model:", metrics.r2_score(y, ada_pred))

Accuracy of the model: 0.8731009548456359
```

The accuracy score is improved to 0.873, which exceeds the target of 0.850.

```
# New data for which predictions are to be made
new_data = pd.DataFrame({'BegInv': [4250], 'EndInv': [4220],
                        'MoveQty': [1862], 'StdCases': [526576]})

# new_data = pd.DataFrame({'BegInv': [12350], 'EndInv': [11940],
#                          'MoveQty': [1312], 'StdCases': [584822]})

# Make predictions on new data
predictions = rf.predict(new_data)

print("Predictions:", predictions)

Predictions: [250284.57195833]
```

## Conclusion :

Based on the analysis of the dataset, four different regression models were used to predict the cashflow variable-payment. The accuracy score for the four models are.

1. Multi-linear regression : 0.109
2. Polynomial regression: 0.774
3. Random forest: 0.826
4. AdaBoost: 0.873

After evaluating the accuracy scores of all four models, it was found that AdaBoost Regression had the highest accuracy score. Therefore, it can be concluded that AdaBoost Regression is the most appropriate model for predicting the fish handling cost in this dataset.